

Neuro-Symbolic = Neural + Logical + Probabilistic

Luc De Raedt¹, Robin Manhaeve¹, Sebastijan Dumančić¹,
Thomas Demeester² and Angelika Kimmig³

¹KU Leuven

²Ghent University - imec

³Cardiff University

{luc.deraedt, robin.manhaeve, sebastijan.dumancic}@cs.kuleuven.be,
thomas.demeester@ugent.be, KimmigA@cardiff.ac.uk

Abstract

The overall goal of neuro-symbolic computation is to integrate high-level reasoning with low-level perception. We argue 1) that neuro-symbolic computation should integrate neural networks with the two most prominent methods for reasoning, that is, logic and probability, and 2) that neuro-symbolic integrated methods should have the pure neural, logical and probabilistic methods *as special cases*. We examine the state-of-the-art with regard to these claims and briefly position our own contribution DeepProbLog in this perspective.

1 Introduction

There is a growing interest in neuro-symbolic computation, that is, in combining high-level reasoning with low-level perception that offers the best of both worlds. The growing consensus that both forms of reasoning are essential to achieve true (artificial) intelligence [Kahneman, 2011] has put the quest for neural symbolic computation [Garcez *et al.*, 2012; Garcez *et al.*, 2015; Hammer and Hitzler, 2007] high on the research agenda. While deep learning excels at low-level perception, there is also a growing awareness of its limitations, certainly in terms of reasoning. Despite various attempts to demonstrate reasoning-like behaviour with deep learning [Santoro *et al.*, 2017], artificial neural networks' current reasoning abilities are nowhere close to what is possible with the typical high-level reasoning approaches. The two most prominent frameworks for reasoning are logic and probability. While in the past, they were studied by separate communities in artificial intelligence, a growing body of researchers is working towards their integration, and even aims at combining probability with logic and statistical learning; cf. the areas of statistical relational artificial intelligence [De Raedt *et al.*, 2016; Getoor and Taskar, 2007] and probabilistic logic programming [De Raedt and Kimmig, 2015]. The reasoning abilities of statistical relational artificial intelligence approaches are complementary to the strong pattern-recognition abilities of deep learning.

2 Position statement

Based on our experience in upgrading machine learning systems towards the use of (probabilistic) logical and relational representations [De Raedt, 2008; Muggleton *et al.*, 2012], we argue that *a first desirable property of frameworks that integrate two other frameworks A and B, is to have the original frameworks A and B as a special case of the integrated one*. If A or B cannot be fully reconstructed, one clearly loses certain abilities, which is not only undesirable but which also implies that there is no true unification. Applying this property to neuro-symbolic computation implies that the existing frameworks should have both the neural and the symbolic representations as special case. When this is the case, one retains both the learning abilities of the neural component as well as the reasoning and learning abilities and the semantics of the symbolic representations. Unfortunately, this property is not satisfied by the vast majority of neuro-symbolic approaches in that they either push the symbolic representation inside the neural network (from which the logic cannot be recovered), or vice versa, apply neural learning principles to symbolic representations (and risk losing both the pure neural component and the logical semantics), as we shall show in the next section.

We also advocate a second desirable property for neuro-symbolic computation: models that learn from observed samples should be able to deal with uncertainty. Therefore, *one should not only integrate logic with neural networks in neuro-symbolic computation, but also probability*.

This effectively leads to an integration of probabilistic logics (hence statistical relational AI) with neural networks and opens up new abilities. Furthermore, although at first sight, this may appear as a complication, it actually can greatly simplify the integration of neural networks with logic. The reason for this is that the probabilistic framework provides a clear optimisation criterion, namely the probability of the training examples. Real-valued probabilistic quantities are also well-suited for gradient-based training procedures, as opposed to discrete logic quantities.

3 State-of-the-art in Neuro-Symbolic Computation

We now examine the state-of-the-art in neuro-symbolic computation with respect to these two properties, i.e., whether

	Neural network / embeddings as special case	Logic as special case	Probabilistic	Underlying logic
Semantic based regularization [Diligenti <i>et al.</i> , 2017]	X			FOL
Logic Tensor Networks [Donadello <i>et al.</i> , 2017]				FOL
Lifted Rule Injection [Demeester <i>et al.</i> , 2016]	X			Implication rules
Adversarial Set Regularisation [Minervini <i>et al.</i> , 2017]	X			Clausal logic
Semantic Loss Function [Xu <i>et al.</i> , 2018]	X		X	Propositional logic
∂ ILP [Evans and Grefenstette, 2018]				Datalog
Neural Theorem Prover [Rocktäschel and Riedel, 2017]				Clausal logic
TensorLog [Cohen <i>et al.</i> , 2018]		X		Datalog Subset
DeepProbLog [Manhaeve <i>et al.</i> , 2018]	X	X	X	Clausal logic

Table 1: Comparing Neuro-Symbolic frameworks

they retain the neural network and logic as special cases, and whether they integrate probabilities. Here, we only consider works that combine neural networks with logic, although the field of neuro-symbolic integration is wider than this. The results are summarized in Table 1.

This subset of the domain can be divided into two categories: logic as constraints and differentiable logic frameworks.



Logic as constraints. The dominant idea is to use logic as a constraint on a deep model. That is, the deep model is extended with a regularization term, derived from the desired logical properties, which encourages it to mimic logical reasoning: not obeying such logical properties induces a penalty. For instance, [Diligenti *et al.*, 2017] and [Donadello *et al.*, 2017] use first-order logic to specify constraints on the output of the neural network, while [Demeester *et al.*, 2016] and [Minervini *et al.*, 2017] use logical IF-THEN rules, derived from expert knowledge, to enforce the embeddings to be more consistent with the logical constraints. [Xu *et al.*, 2018] introduce a generalization of this principle such that more complex logical constraints can be imposed on any deep model. Common to all these approaches is the use of logic as a regularizer, which encourages deep models to satisfy the constraints. This leads to the logic being encoded into the parameters (either into the weights of the neural network, or directly into the embeddings). The constraints are soft, and there is no guarantee that they all will be satisfied. In fact, the trade-off between encouraging the constraints vs. following the data in case both are contradictory, is merely a hyperparameter of the model. From a logical perspective, constraints only support one form of reasoning, an alternative form of reasoning is based on providing definitions of predicates (in the form of rules or definite clauses), and to use such rules to answer queries. This is the basis for the popular programming and database languages Prolog and Datalog. The aforementioned methods thus generally do have neural networks as a special case (i.e., when there are no additional constraints), however, they do not have logic as a special case when leaving out the neural part. The importance of being able to fully recover the logic is already hinted at in some of these papers. For example, [Xu *et al.*, 2018] show that neural networks trained with the additional regularization term cannot consistently make predictions coherent with the logic they were trained on.

Differentiable logic frameworks. A different class of neuro-symbolic systems works by making the logic program

differentiable, which is achieved by reformulating the basic reasoning primitives using the mathematics of differentiable functions. [Rocktäschel and Riedel, 2017] apply this to the backward reasoning procedure of Prolog, while [Evans and Grefenstette, 2018] do this in the style of forward reasoning. Although both systems are based on standard reasoning algorithms, the original logic cannot be recovered as special case, since the original semantics and inference have been fundamentally changed. Similarly, [Cohen *et al.*, 2018] introduce a framework to compile a tractable subset of logic programs into differentiable functions and to execute it with neural networks. All three systems mentioned above construct a neural network to perform the execution, but they do not have neural networks as a special case. These approaches can only deal with rather restricted classes of logic programs. Furthermore, because they push the logic programming components into the neural networks, both the proper semantics of logic programs and the programming language nature of the framework are sacrificed. This is clear when considering that the approach of [Rocktäschel and Riedel, 2017] is to map a proof tree on a neural network, which can then be differentiated as the logic has been removed. These systems also suffer from not keeping the original logic as a special case. For example, [Rocktäschel and Riedel, 2017] and [Evans and Grefenstette, 2018] both mention that the neural execution of the differentiable logic creates a large computational overhead.

4 Our approach: DeepProbLog

The proposed approach is pursued in DeepProbLog [Manhaeve *et al.*, 2018], a neuro-symbolic framework developed in the past year. What sets DeepProbLog apart from the other approaches mentioned above is that instead of changing the logic, it is extended. To create DeepProbLog, we extended the probabilistic logic programming language ProbLog [Fierens *et al.*, 2015] with neural predicates. The underlying concept of a probabilistic logic programming language is simple: (ground) atomic expressions of the form $q(t_1, \dots, t_n)$ (aka tuples in a relational database) are considered as (independent) random variables that have a probability p of being true. This view, which was pioneered in the distribution semantics by [Sato, 1995], effectively unifies the basic primitives in logic with those in probability theory: propositions become random variables. On top of these “probabilistic facts” one can then define rules (in the form of logic programs) that allow to derive conclusions from these facts

and induce a probability distribution over possible worlds. The neural predicate extends this concept so that atomic expressions can also be annotated with the output of a neural network, given that the output can be considered a probability. This simple idea is appealing as it allows us to retain all the essential components of the probabilistic logic programming language: the semantics, the inference mechanism, as well as the implementation. At the same time, it clearly decouples the logical component from the neural component. As neural networks become predicates, they can be called from the logic programs and so the logical layer is at the higher-level and the neural one at the lower-level with the neural predicates providing the interface, while both sides can treat each other as a black box. For example, consider the task where we have to learn to classify sums from pairs of handwritten digits, e.g.  +  = 8. In line with our position statement, the neural network needs to handle *only* the image recognition, while the addition is handled purely by logic. This task can thus easily be solved in DeepProbLog by specifying a neural predicate that defines the classification of single digits, and a single line of logic that defines the addition. As shown in [Manhaeve *et al.*, 2018], this single line of logic is enough to clearly outperform a fully neural baseline, as it converges quicker, to a higher accuracy. Furthermore, because the neural network in the DeepProbLog program classifies digits, it can be reused for other tasks after training, whereas the baseline can only be used on the same task. It is easy to see how both the logic and the neural networks are preserved as special cases: if the DeepProbLog program has no neural predicates, it simply becomes a normal ProbLog program. If that program has no probabilistic facts, it becomes a pure logic program. Furthermore, if the program only contains a neural predicate, it becomes equivalent to directly training the neural network. In describing the learning algorithm for DeepProbLog it is important to note that DeepProbLog, in contrast to other approaches for neuro-symbolic learning, not only encompasses a neural and a logical framework, but also a probabilistic one. Making the logic probabilistic also makes it differentiable, allowing gradients to be derived so that the neural networks can be trained with standard gradient descent methods. For the given example, note that the neural predicate is trained from an indirect signal, i.e., from observations of sums (at the output of the logical layer), rather than through labels of individual digits.

To summarize, DeepProbLog occupies a unique position in the space of neuro-symbolic methods in that i) it is a programming language that supports neural networks and machine learning, and it has a well-defined semantics (as an extension of Prolog, it is Turing equivalent); (ii) it integrates logical reasoning with neural networks; so both symbolic and subsymbolic representations and inference, learning and reasoning, perception and inference; (iii) it integrates probabilistic modeling, programming and reasoning with neural networks (as DeepProbLog extends the probabilistic programming language ProbLog, which can be regarded as a very expressive directed graphical modeling language [De Raedt *et al.*, 2016]).

Acknowledgements

This work has been partially supported by the European Research Council Advanced Grant project SYNTH9 (ERCAdG-694980). RM is a SB PhD fellow at FWO (1S61718N).

References

- [Cohen *et al.*, 2018] William W Cohen, Fan Yang, and Kathryn Rivard Mazaitis. Tensorlog: Deep learning meets probabilistic databases. *Journal of Artificial Intelligence Research*, 1:1–15, 2018.
- [De Raedt and Kimmig, 2015] Luc De Raedt and Angelika Kimmig. Probabilistic (logic) programming concepts. *Machine Learning*, 100(1):1–43, 2015.
- [De Raedt *et al.*, 2016] Luc De Raedt, Kristian Kersting, Sri-raam Natarajan, and David Poole. Statistical relational artificial intelligence: Logic, probability and computation. *Synthesis Lectures on Artificial Intelligence and Machine Learning*, 2016.
- [De Raedt, 2008] Luc De Raedt. *Logical and Relational Learning*. Springer, 2008.
- [Demeester *et al.*, 2016] Thomas Demeester, T. Rocktäschel, and S. Riedel. Lifted rule injection for relation embeddings. In *EMNLP*, 2016.
- [Diligenti *et al.*, 2017] Michelangelo Diligenti, Marco Gori, and Claudio Sacca. Semantic-based regularization for learning and inference. *Artificial Intelligence*, 244:143–165, 2017.
- [Donadello *et al.*, 2017] Ivan Donadello, Luciano Serafini, and Artur S. d’Avila Garcez. Logic tensor networks for semantic image interpretation. In *IJCAI*, 2017.
- [Evans and Grefenstette, 2018] Richard Evans and Edward Grefenstette. Learning explanatory rules from noisy data. *JAIR*, 2018.
- [Fierens *et al.*, 2015] Daan Fierens, Guy Van den Broeck, Joris Renkens, Dimitar Shterionov, Bernd Gutmann, Ingo Thon, Gerda Janssens, and Luc De Raedt. Inference and learning in probabilistic logic programs using weighted Boolean formulas. *Theory and Practice of Logic Programming*, 15(3):358–401, 2015.
- [Garcez *et al.*, 2012] Artur S d’Avila Garcez, Krysia B Broda, and Dov M Gabbay. *Neural-symbolic learning systems: foundations and applications*. Springer Science & Business Media, 2012.
- [Garcez *et al.*, 2015] Artur d’Avila Garcez, Tarek R Besold, Luc De Raedt, Peter Földiák, Pascal Hitzler, Thomas Icard, Kai-Uwe Kühnberger, Luis C Lamb, Risto Miikkulainen, and Daniel L Silver. Neural-symbolic learning and reasoning: contributions and challenges. In *2015 AAAI Spring Symposium Series*, 2015.
- [Getoor and Taskar, 2007] Lise Getoor and Ben Taskar. *Introduction to statistical relational learning*. MIT press, 2007.
- [Hammer and Hitzler, 2007] Barbara Hammer and Pascal Hitzler. *Perspectives of neural-symbolic integration*, volume 8. Springer Heidelberg, 2007.

- [Kahneman, 2011] Daniel Kahneman. *Thinking, fast and slow*. Farrar, Straus and Giroux New York, 2011.
- [Manhaeve *et al.*, 2018] Robin Manhaeve, Sebastijan Dumančić, Angelika Kimmig, Thomas Demeester, and Luc De Raedt. Deepproblog: Neural probabilistic logic programming. In *NeurIPS*, 2018.
- [Minervini *et al.*, 2017] Pasquale Minervini, Thomas Demeester, Tim Rocktäschel, and Sebastian Riedel. Adversarial sets for regularising neural link predictors. In *UAI*, 2017.
- [Muggleton *et al.*, 2012] Stephen Muggleton, Luc De Raedt, David Poole, Ivan Bratko, Peter Flach, Katsumi Inoue, and Ashwin Srinivasan. Ilp turns 20. *Machine learning*, 86(1):3–23, 2012.
- [Rocktäschel and Riedel, 2017] Tim Rocktäschel and Sebastian Riedel. End-to-end differentiable proving. In *NIPS*, 2017.
- [Santoro *et al.*, 2017] Adam Santoro, David Raposo, David G Barrett, Mateusz Malinowski, Razvan Pascanu, Peter Battaglia, and Tim Lillicrap. A simple neural network module for relational reasoning. In *NIPS*, 2017.
- [Sato, 1995] T. Sato. A statistical learning method for logic programs with distribution semantics. In *ICLP*, 1995.
- [Xu *et al.*, 2018] Jingyi Xu, Zilu Zhang, Tal Friedman, Yitao Liang, and Guy Van den Broeck. A semantic loss function for deep learning with symbolic knowledge. In *ICML*, 2018.

Made with the new Google Sites [Create a site](#) [Report abuse](#)

This site uses cookies from Google to deliver its services and to analyze traffic. Information about your use of this site is shared with Google. By using this site, you agree to its use of cookies.

[LEARN MORE](#) [GOT IT](#)



S
t
a
t
i
s
t
i
c
a
l
m
a
c
h
i
n
e
l
e
a
r
n
i
n
g
a

n
d
c
o
n
n
e
c
t
i
o
n
i
s
t
s
y
s
t
e
m
s
h
a
v
e
a
c
h
i
e
v
e
d
i
n
d
u
s
t

r
i
a
l
r
e
l
e
v
a
n
c
e
i
n
a
n
u
m
b
e
r
o
f
a
r
e
a
s
,
f
r
o
m
r
e
t
a
i
l

t
o
h
e
a
l
t
h
,
b
y
t
h
e
i
r
s
t
a
t
e
-
o
f
-
t
h
e
-
a
r
t
p
e
r
f
o
r
m
a

n
c
e
i
n
l
a
n
g
u
a
g
e
m
o
d
e
l
l
i
n
g
,
s
p
e
e
c
h
r
e
c
o
g
n
i
t
i
o
n

,
g
r
a
p
h
a
n
a
l
y
t
i
c
s
,
i
m
a
g
e
,
v
i
d
e
o
a
n
d
s
e
n
s
o
r
d
a
t
a

a
n
a
l
y
s
i
s
.
S
y
m
b
o
l
i
c
s
y
s
t
e
m
s
a
r
e
c
h
a
l
l
e
n
g
e
d
b
y
u

n
s
t
r
u
c
t
u
r
e
d
d
a
t
a
,
b
u
t
t
h
e
y
a
r
e
r
e
c
o
g
n
i
z
e
d
a
s
b
e

i
n
g
i
n
p
r
i
n
c
i
p
l
e
t
r
a
n
s
p
a
r
e
n
t
,
i
n
t
h
a
t
r
e
a
s
o
n
e
d

f
a
c
t
s
f
r
o
m
k
n
o
w
l
e
d
g
e
b
a
s
e
s
c
a
n
b
e
i
n
s
p
e
c
t
e
d
t
o
i

n
t
e
r
p
r
e
t
h
o
w
d
e
c
i
s
i
o
n
s
f
o
l
l
o
w
f
r
o
m
i
n
p
u
t
.
C
o
n
n

e
c
t
i
o
n
i
s
t
a
n
d
s
y
m
b
o
l
i
c
m
e
t
h
o
d
s
a
l
s
o
c
o
n
t
r
a
s
t
i

n
t
h
e
p
r
o
b
l
e
m
s
t
h
e
y
e
x
c
e
l
a
t
:
s
c
e
n
e
r
e
c
o
g
n
i
t
i
o
n

f
r
o
m
i
m
a
g
e
s
a
p
p
e
a
r
s
t
o
b
e
a
p
r
o
b
l
e
m
s
t
i
l
l
o
u
t
s
i
d

e
t
h
e
c
a
p
a
b
i
l
i
t
i
e
s
o
f
s
y
m
b
o
l
i
c
s
y
s
t
e
m
s
f
o
r
e
x
a
m

p
l
e
,
w
h
i
l
e
n
e
u
r
a
l
n
e
t
s
a
r
e
n
o
t
y
e
t
s
u
f
f
i
c
i
e
n
t
f
o

r
i
n
d
u
s
t
r
i
a
l
-
s
t
r
e
n
g
t
h
c
o
m
p
l
e
x
p
l
a
n
n
i
n
g
s
c
e
n
a

r
i
o
s
a
n
d
d
e
d
u
c
t
i
v
e
r
e
a
s
o
n
i
n
g
t
a
s
k
s
.
N
e
u
r
a
l
-
S
y

m
b
o
l
i
c
c
o
m
p
u
t
a
t
i
o
n
a
i
m
s
t
o
b
u
i
l
d
r
i
c
h
c
o
m
p
u
t
a
t

a
i
o
n
a
l
m
o
d
e
l
s
a
n
d
s
y
s
t
e
m
s
b
y
c
o
m
b
i
n
i
n
g
b
o
t
h
c
o
n

n
e
c
t
i
o
n
i
s
t
a
n
d
s
y
m
b
o
l
i
c
l
e
a
r
n
i
n
g
p
a
r
a
d
i
g
m
s
.
T

h
i
s
c
o
m
b
i
n
a
t
i
o
n
h
o
p
e
s
t
o
f
o
r
m
s
y
n
e
r
g
i
e
s
b
e
t
w
e
e

n
t
h
e
i
r
s
t
r
e
n
g
t
h
s
w
h
i
l
e
o
v
e
r
c
o
m
i
n
g
t
h
e
i
r
c
o
m
p
l

e
m
e
n
t
a
r
y
w
e
a
k
n
e
s
s
e
s
.
T
h
e
N
e
S
y
w
o
r
k
s
h
o
p
s
e
r
i
e
s

i
s
s
e
e
n
a
s
a
m
a
j
o
r
v
e
n
u
e
f
o
r
t
h
e
p
r
e
s
e
n
t
a
t
i
o
n
a
n
d

d
i
s
c
u
s
s
i
o
n
o
f
k
e
y
t
o
p
i
c
s
r
e
l
a
t
e
d
t
o
n
e
u
r
a
l
-
s
y
m

b
o
l
i
c
c
o
m
p
u
t
i
n
g
.
N
e
S
y
h
a
s
p
r
o
v
i
d
e
d
a
n
a
t
m
o
s
p
h
e

r
e
f
o
r
t
h
e
f
r
e
e
e
x
c
h
a
n
g
e
o
f
i
d
e
a
s
a
n
d
n
e
t
w
o
r
k
i
n
g

w
i
t
h
i
n
t
h
e
c
o
m
m
u
n
i
t
y
o
f
s
c
i
e
n
t
i
s
t
s
t
h
a
t
s
t
r
a
d

I
e
t
h
e
l
i
n
e
b
e
t
w
e
e
n
s
t
a
t
i
s
t
i
c
a
l
a
n
d
s
y
m
b
o
l
i
c
m
e

t
h
o
d
s
i
n
A
I
.

- T
a
r
e
k
B
e
s
- A
t
t
u
A
b
p
A
a
H
e
a
G
t
p
A
e
L

a
,
e
i
T
y
!
0
h
p
v
e
P
▪ s
G
e
o
o
g
e
C
y
h
e
n
k
u
K
▪ D
P
f
e
m
a
y
L
e
▪ c
D
e
r

le
N
B
b
A
a
n
d
W
h
a
g
b
s
S
M
a
n
e
U
e
a
v
e
C
a
n
a
y
■ a
■ A
D
e
u
e
k
D
A
v
a

h
a
~~W~~
a
i
g
b
z
S
C
a
t
~~g~~
U
n
i
v
e
r
s
i
t
y
o
f
B
A
n
d
o
n
.
F
r
a
n
V
a
n
H

a
r
m
e
l
e
n
,
V
r
i
j
e
U
n
i
v
e
r
s
i
t
e
i
t
A
m
s
t
e
r
d
a
m

■

l
r
i
n
a

H
i
g
g
i
n
s
,
G
o
o
g
l
e

▪

P
a
s
c
a
l
H
i
t
z
l
e
r
,
W
r
i
g
h
t
S
t
a
t
e

U
n
i
v
e
r
s
i
t
y

▪

K
r
i
s
t
a
n
K
e
r
s
t
i
n
g
,
T
U
D
a
r
m
s
t
a
d
t

▪

L

u
i
s
L
a
m
b
,
F
e
d
e
r
a
l
U
n
i
v
e
r
s
i
t
y
o
f
R
i
o
G
r
a
n
d
e
d
o
S
u

I

■

F
r
e
d
d
y
L
e
c
u
e
,
I
N
R
I
A
a
n
d
T
h
a
l
e
s

■

T
h
o
m
a
s
L
u
k
a
s

i
e
w
i
c
z
,
U
n
i
v
e
r
s
i
t
y
o
f
O
x
f
o
r
d

■

P
a
s
q
u
a
l
e
M
i
n
e
r
v

i
n
i
,
U
n
i
v
e
r
s
i
t
y
C
o
l
l
e
g
e
L
o
n
d
o
n

■

L
u
c
D
e
R
a
e
d
t
,
K

a
t
h
o
l
i
e
k
e
U
n
i
v
e
r
s
i
t
e
i
t
L
e
u
v
e
n

■

L
u
c
i
a
n
o
S
e
r
a
f

i
n
i
,
F
o
n
d
a
z
i
o
n
e
B
r
u
n
o
K
e
s
s
l
e
r

■

A
m
i
n
a
S
h
a
b
b
e
e
r

,
A
m
a
z
o
n

▪

T
r
a
n
S
o
n
,
U
n
i
v
e
r
s
i
t
y
o
f
T
a
s
m
a
n
i
a

▪

M
i
c

h
a
e
l
S
p
r
a
n
g
e
r
,
S
o
n
y
C
o
m
p
u
t
e
r
S
c
i
e
n
c
e
L
a
b
s

C
a

I
I
F
o
r
P
a
p
e
r
s

N
e
S
y
i
n
v
i
t
e
s
t
h
e
o
r
e
t
i
c
a
l
a
n
d

a
p
p
l
i
e
d
s
u
b
m
i
s
s
i
o
n
s
o
f
a
l
l
t
y
p
e
s
t
h
a
t
s
p
a
n
b
o
t
h

c
o
n
n
e
c
t
i
o
n
i
s
t
a
n
d
s
y
m
b
o
l
i
c
l
e
a
r
n
i
n
g
p
a
r
a
d
i
g
m

s
.
W
e
f
u
r
t
h
e
r
i
n
v
i
t
e
p
a
p
e
r
s
d
e
t
a
i
l
i
n
g
e
x
p
e
r
i
m
e

n
t
a
l
a
n
d
i
n
-
t
h
e
-
w
i
l
d
n
e
u
r
a
l
-
s
y
m
b
o
l
i
c
s
y
s
t
e
m
s

a
n
d
p
a
p
e
r
s
o
n
t
o
p
i
c
s
w
h
e
r
e
n
e
u
r
a
l
-
s
y
m
b
o
l
i
c
l
e
a

r
n
i
n
g
h
a
s
a
s
t
r
o
n
g
u
s
e
c
a
s
e
.
T
o
p
i
c
s
o
f
i
n
t
e
r
e
s
t
i

n
c
l
u
d
e
,
b
u
t
a
r
e
n
o
t
l
i
m
i
t
e
d
t
o
,
t
h
e
f
o
l
l
o
w
i
n
g
:
■

K
n
o
w
l
e
d
g
e
r
e
p
r
e
s
e
n
t
a
t
i
o
n
a
n
d
r
e
a
s
o
n
i
n
g
i
n
(
d
e

e
p
)
n
e
u
r
a
l
n
e
t
w
o
r
k
s

■

S
y
m
b
o
l
i
c
k
n
o
w
l
e
d
g
e
e
x
t
r
a

c
t
i
o
n
f
r
o
m
n
e
u
r
a
l
a
n
d
s
t
a
t
i
s
t
i
c
a
l
l
e
a
r
n
i
n
g
m
o
d

e
l
s
.
E
x
p
l
a
i
n
a
b
l
e
A
l
m
o
d
e
l
s
,
s
y
s
t
e
m
s
,
a
n
d
t
e
c
h
n

i
q
u
e
s
t
h
a
t
i
n
t
e
g
r
a
t
e
c
o
n
n
e
c
t
i
o
n
i
s
t
a
n
d
s
y
m
b
o
l

i
c
p
a
r
a
d
i
g
m
s

■

N
e
u
r
a
l
-
s
y
m
b
o
l
i
c
c
o
g
n
i
t
i
v
e
m
o
d
e

l
s
■
B
i
o
l
o
g
i
c
a
l
l
y
-
i
n
s
p
i
r
e
d
n
e
u
r
a
l
-
s
y
m
b
o
l
i
c
i

n
t
e
g
r
a
t
i
o
n

■

C
o
n
t
i
n
u
a
l
l
e
a
r
n
i
n
g
,
i
n
t
e
g
r
a
t
i
o
n

o
f
l
o
g
i
c
a
n
d
p
r
o
b
a
b
i
l
i
t
i
e
s
w
i
t
h
n
e
u
r
a
l
n
e
t
w
o
r
k

s

■

N
e
u
r
a
l
-
s
y
m
b
o
l
i
c
m
e
t
h
o
d
s
f
o
r
s
t
r
u
c
t
u
r
e
d
l
e
a

r
n
i
n
g
t
a
s
k
s
,
i
n
c
l
u
d
i
n
g
t
r
a
n
s
f
e
r
,
m
e
t
a
,
a
n
d
r
e
l

a
t
i
o
n
a
l
l
e
a
r
n
i
n
g

▪

N
o
v
e
l
c
o
n
n
e
c
t
i
o
n
i
s
t
s
y
s
t
e
m

s
a
b
l
e
t
o
p
e
r
f
o
r
m
t
r
a
d
i
t
i
o
n
a
l
l
y
s
y
m
b
o
l
i
c
A
l
t
a
s

k
s
(
e
.
g
.
,
l
o
g
i
c
a
l
d
e
d
u
c
t
i
o
n
)

■

N
o
v
e
l
s
y
m
b
o
l
i
c
s

y
s
t
e
m
s
a
b
l
e
t
o
p
e
r
f
o
r
m
t
r
a
d
i
t
i
o
n
a
l
l
y
c
o
n
n
e
c
t
i

o
n
i
s
t
t
a
s
k
s
(
e
.
g
.
,
u
n
s
t
r
u
c
t
u
r
e
d
d
a
t
a
a
n
a
l
y
s
i
s

)

■

A
p
p
l
i
c
a
t
i
o
n
s
i
n
s
i
m
u
l
a
t
i
o
n
,
f
i
n
a
n
c
e
,
r
o
b
o
t

i
c
s
,
t
h
e
s
e
m
a
n
t
i
c
w
e
b
,
s
o
f
t
w
a
r
e
e
n
g
i
n
e
e
r
i
n
g
,
s

y
s
t
e
m
s
e
n
g
i
n
e
e
r
i
n
g
,
b
i
o
i
n
f
o
r
m
a
t
i
c
s
,
a
n
d
v
i
s
u

a
l
i
n
t
e
l
l
i
g
e
n
c
e
.

N
e
S
y
a
l
s
o
i
n
v
i
t
e
s
s
h
o
r
t
,
c
l
e
a

r
,
a
n
d
w
e
l
l
-
w
r
i
t
t
e
n
p
o
s
i
t
i
o
n
p
a
p
e
r
s
f
o
r
p
r
e
s
e
n

t
a
t
i
o
n
a
n
d
d
i
s
c
u
s
s
i
o
n
d
u
r
i
n
g
t
h
e
w
o
r
k
s
h
o
p
.

S
u

b
m
i
s
s
i
o
n

R
e
s
e
a
r
c
h
e
r
s
a
n
d
p
r
a
c
t
i
t
i
o
n
e
r
s
a
r

e
i
n
v
i
t
e
d
t
o
s
u
b
m
i
t
o
r
i
g
i
n
a
l
p
a
p
e
r
s
u
s
i
n
g
t
h
e
I
J

C
A
l
p
a
p
e
r
f
o
r
m
a
t
t
h
a
t
h
a
v
e
n
o
t
b
e
e
n
s
u
b
m
i
t
t
e
d
f
o

r
r
e
v
i
e
w
o
r
p
u
b
l
i
s
h
e
d
e
l
s
e
w
h
e
r
e
.
S
u
b
m
i
t
t
e
d
p
a
p

e
r
s
m
u
s
t
b
e
w
r
i
t
t
e
n
i
n
E
n
g
l
i
s
h
a
n
d
s
h
o
u
l
d
n
o
t
e
x
c

e
e
d
6
p
a
g
e
s
i
n
t
h
e
c
a
s
e
o
f
r
e
s
e
a
r
c
h
a
n
d
e
x
p
e
r
i
e
n
c

e
p
a
p
e
r
s
,
o
r
3
p
a
g
e
s
i
n
t
h
e
c
a
s
e
o
f
p
o
s
i
t
i
o
n
p
a
p
e
r

s
(
i
n
c
l
u
d
i
n
g
f
i
g
u
r
e
s
,
b
i
b
l
i
o
g
r
a
p
h
y
a
n
d
a
p
p
e
n
d

i
c
e
s
)
.
A
l
l
s
u
b
m
i
t
t
e
d
p
a
p
e
r
s
w
i
l
l
b
e
j
u
d
g
e
d
b
a
s
e

d
o
n
t
h
e
i
r
q
u
a
l
i
t
y
,
r
e
l
e
v
a
n
c
e
,
o
r
i
g
i
n
a
l
i
t
y
,
s
i

g
n
i
f
i
c
a
n
c
e
,
a
n
d
s
o
u
n
d
n
e
s
s
.
S
u
b
m
i
s
s
i
o
n
s
e
d
n

o
t
b
e
d
o
u
b
l
e
-
b
l
i
n
d
.
S
u
b
m
i
t
y
o
u
r
p
a
p
e
r
v
i
a
:
[E](#)
[a](#)
[s](#)
[y](#)

[C](#)
[h](#)
[a](#)
[i](#)
[r](#)
[S](#)
[u](#)
[b](#)
[m](#)
[i](#)
[s](#)
[s](#)
[i](#)
[o](#)
[n](#)
[L](#)
[i](#)
[n](#)
[k](#)

P
r
e
s
e
n
t
a
t
i
o
n

S
e
l
e

c
t
e
d
p
a
p
e
r
s
w
i
l
l
b
e
p
r
e
s
e
n
t
e
d
d
u
r
i
n
g
t
h
e
w
o
r
k
s
h

o
p
.
T
h
e
w
o
r
k
s
h
o
p
w
i
l
l
i
n
c
l
u
d
e
e
x
t
r
a
t
i
m
e
f
o
r
a
u
d

i
e
n
c
e
d
i
s
c
u
s
s
i
o
n
o
f
t
h
e
p
r
e
s
e
n
t
a
t
i
o
n
a
l
l
o
w
i
n
g

t
h
e
g
r
o
u
p
t
o
h
a
v
e
a
b
e
t
t
e
r
u
n
d
e
r
s
t
a
n
d
i
n
g
o
f
t
h
e
i

s
s
u
e
s
,
c
h
a
l
l
e
n
g
e
s
,
a
n
d
i
d
e
a
s
b
e
i
n
g
p
r
e
s
e
n
t
e
d
.

A
p
r
e
s
e
n
t
a
t
i
o
n
m
a
y
s
e
l
e
c
t
e
d
f
o
r
o
r
a
l
o
r
p
o
s
t
e
r
p

r
e
s
e
n
t
a
t
i
o
n
d
e
p
e
n
d
i
n
g
o
n
t
h
e
n
u
m
b
e
r
o
f
a
c
c
e
p
t
e

d
p
a
p
e
r
s
a
n
d
o
n
t
h
e
i
d
e
a
l
p
r
e
s
e
n
t
a
t
i
o
n
s
t
y
l
e
f
o
r

t
h
e
s
u
b
m
i
t
t
e
d
w
o
r
k
.

P
u
b
l
i
c
a
t
i
o
n

A
c
c
e
p
t
e

d
p
a
p
e
r
s
r
e
g
a
r
d
l
e
s
s
o
f
p
r
e
s
e
n
t
a
t
i
o
n
f
o
r
m
a
t
w
i
l

l
b
e
p
u
b
l
i
s
h
e
d
i
n
o
f
f
i
c
i
a
l
w
o
r
k
s
h
o
p
p
r
o
c
e
e
d
i
n
g

s
,
w
h
i
c
h
w
i
l
l
b
e
d
i
s
t
r
i
b
u
t
e
d
d
u
r
i
n
g
t
h
e
w
o
r
k
s
h
o

p
.
A
u
t
h
o
r
s
o
f
p
a
p
e
r
s
t
h
a
t
a
r
e
w
e
l
l
-
r
e
v
i
e
w
e
d
o
r
i

n
c
i
t
e
s
i
n
t
e
r
e
s
t
d
u
r
i
n
g
t
h
e
w
o
r
k
s
h
o
p
w
i
l
l
b
e
i
n
v

i
t
e
d
t
o
s
u
b
m
i
t
a
r
e
v
i
s
e
d
a
n
d
e
x
t
e
n
d
e
d
v
e
r
s
i
o
n
o
f

t
h
e
i
r
p
a
p
e
r
s
t
o
t
h
e
J
o
u
r
n
a
l
o
f
A
p
p
l
i
e
d
L
o
g
i
c
s
,
C

o
l
l
e
g
e
P
u
b
l
i
c
a
t
i
o
n
s
.